

# Spark RDD-based programming

February 22, 2024

## 1 Analyze bike-sharing system of Barcelona - Spark RDD-based programming

```
[ ]: # In this analysis, I am going to consider the occupancy of the stations where  
      ↪users can pick up or drop off  
      # bikes in order to identify the most "critical" timeslots (day of the week,  
      ↪hour) for each station.
```

```
[ ]: # data is located on the big data cluster and I am going to read data from  
      ↪there.  
      # there are two types of data:  
      # 1. register.csv: This contains the historical information about the number of  
      ↪used and free slots for  
      # ~3000 stations from May 2008 to September 2008. Each line of register.csv  
      # corresponds to one reading about the situation of one station at a  
      ↪specific timestamp.  
      # 2. stations.csv: It contains the description of the stations (station_id,  
      ↪latitude, longitude, name).
```

```
[1]: registerPath = "/data/students/bigdata_internet/lab3/register.csv"  
      stationPath = "/data/students/bigdata_internet/lab3/stations.csv"
```

```
[ ]: '''In this analysis, PySpark was utilized for its robust distributed computing  
      ↪capabilities,  
      ideal for handling large datasets efficiently.  
      If you're using the PySpark shell, no additional setup is necessary.  
      However, for those working in a Python environment, setting up PySpark involves  
      ↪the following steps:  
      1. Install PySpark: Begin by installing PySpark using pip:  
      pip install pyspark  
      2. Configure PySpark: In your Python script or interactive session, include the  
      ↪following configuration  
      to initialize PySpark:  
      ```python  
      from pyspark import SparkConf, SparkContext  
      conf = SparkConf().setAppName("MyApp")
```

```

sc = SparkContext(conf=conf)
'''
Ensure to execute this configuration before performing any PySpark operations.
For comprehensive installation and configuration instructions, refer to the
    ↪official PySpark documentation:
PySpark Installation Guide
'''

```

```
[ ]: # Reading register data as a RDD
```

```
[2]: registerRDD = sc.textFile(registerPath)
```

```
[ ]: # The file is separated with tab so I will split each row by \t.
```

```
[3]: registerRDDList = registerRDD.map(lambda l: l.split('\t'))
```

```
[5]: registerRDD.count()
```

```
[5]: 25319029
```

```
[ ]: # To clean the data, I am going to filter data that their used_slot != 0 or
    ↪their free_slot != 0
    # because whether there are some bicycles in station or not and it is not
    ↪possible to have 0 for both.
```

```
[4]: registerRDDFiltered = registerRDDList.filter(lambda l: l[2] != "0" or l[3] !=
    ↪"0")
```

```
[7]: registerRDDFiltered.count()
```

```
[7]: 25104122
```

```
[ ]: # There are 25,319,029 rows (one row is for the header) in the original file
    ↪and
    # it decreases to 25,104,122 (one row for the header) after we did the filter
    ↪and deleted wrong data.
```

.....

```
[ ]: # Reading station data as a RDD
```

```
[5]: stationRDD = sc.textFile(stationPath)
```

```

[6]: stationRDDList = stationRDD.map(lambda l: l.split('\t'))

[ ]: # Write a Spark application that selects the pairs (station, timeslot) that are
    ↳ characterized
    # by a high "criticality" value

[ ]: # In this section I am going to find critical stations which have the most used
    ↳ bicycles

[ ]: # Because the file is csv, there is header that I have to remove it because I
    ↳ will analyze by RDD-based programming

[7]: headerR = registerRDDFiltered.first()

[8]: registerCleanRDD = registerRDDFiltered.filter(lambda x: x != headerR)

[9]: headerS = stationRDDList.first()

[10]: stationCleanRDD = stationRDDList.filter(lambda x: x != headerS)

[ ]: # For this analysis I will use "day of week" and "hour" to find critical
    ↳ timeslots
    # So, I am changing the timestamp into this format

[11]: from datetime import datetime as dt

[12]: def format_timestamp(l):
        timestamp = dt.strptime(l[1], "%Y-%m-%d %H:%M:%S")
        formatted_timestamp = dt.strftime(timestamp, "%A, %H")
        l[1] = formatted_timestamp
        return l

[13]: registerTimeslot = registerCleanRDD.map(format_timestamp)

.....

[ ]: # Computes the criticality value  $C(S_i, T_j)$  for each pair  $(S_i, T_j)$ 

[ ]: # Turn register data into  $(k, v)$  pairs of  $(station\_id, timeslot)$  and  $(used\_slot,$ 
    ↳  $free\_slot)$ 

[14]: registerKeyValue = registerTimeslot.map(lambda l: ((l[0], l[1]), [l[2], l[3]]))

[ ]: # Filter only those data that have free_slot = 0 which means that all of the
    ↳ slots were used

```

```

[15]: zeroFreeSlots = registerKeyValue.filter(lambda t: t if t[1][1] == '0' else None)

[ ]: # Turn this data into (k, v) pairs of (station_id, timeslot) and 1 in order to
    ↪ be able to find the number of
    # (station_id, timeslot) with zero free_slot (all bicycles were used)

[16]: zeroNumber = zeroFreeSlots.map(lambda x: (x[0], 1))

[17]: numberZero = zeroNumber.reduceByKey(lambda a,b: a+b)

[ ]: # Turn the register data into (k, v) pairs of (station_id, timeslot) and 1 in
    ↪ order to be able to find
    # the number of all pairs (station_id, timeslot) readings.

[18]: registerNumber = registerKeyValue.map(lambda l: (l[0], 1))

[19]: numberTotal = registerNumber.reduceByKey(lambda a,b: a+b)

[ ]: # Join two previous data (number of free_slots = 0 and all readings) for each
    ↪ pair (station_id, timeslot)

[20]: joinedZeroTotal = numberZero.join(numberTotal)

[ ]: # The ration between these two data will give us the criticality value of each
    ↪ pair (station_id, timeslot)

[21]: criticalityRDD = joinedZeroTotal.map(lambda l: (l[0], int(l[1][0])/
    ↪ int(l[1][1])))

[ ]: # Now, I will select only the critical pairs (Si, Tj) having a criticality
    ↪ value C(Si, Tj) greater than
    # a minimum threshold (0.6).

[22]: criticalPointsRDD = criticalityRDD.filter(lambda x: float(x[1])>=0.6)

[ ]: # Order the results by increasing criticality.

[23]: orderedCriticalPointsRDD = criticalPointsRDD.sortBy(lambda x: float(x[1]), True)

[ ]: # Show the most critical (station_id, timeslot) in Barcelona

[25]: orderedCriticalPointsRDD.collect()

```

```
[25]: [((('9', 'Friday, 10'), 0.6129032258064516),
      (('10', 'Saturday, 00'), 0.622107969151671),
      (('58', 'Monday, 01'), 0.6239554317548747),
      (('9', 'Friday, 22'), 0.6258389261744967),
      (('58', 'Monday, 00'), 0.6323119777158774)]
```

---

```
[ ]: # Store the sorted critical pairs  $C(S_i, T_j)$  in the output folder (also an
      ↪ argument of the application),
      # by using a csv files (with header), where columns are separated by "tab".
      ↪ Store exactly the following
      # attributes separated by a "tab":
      # station / station longitude / station latitude / day of week / hour /
      ↪ criticality value
```

```
[24]: orderedCriticalSeparated = orderedCriticalPointsRDD.map(lambda x: (x[0][0],
      ↪ (x[0][1], x[1])))
```

```
[25]: stationPairRDD = stationCleanRDD.map(lambda s: (s[0], (s[1], s[2])))
```

```
[ ]: # Join critical stations RDD from register data with station data
```

```
[26]: joinedCriticalStationsRDD = stationPairRDD.join(orderedCriticalSeparated)
```

```
[27]: finalRDD = joinedCriticalStationsRDD.map(lambda s: [s[0], s[1][0][0],
      ↪ s[1][0][1], s[1][1][0].split(',')[0], s[1][1][0].split(',')[1], s[1][1][1]])
```

```
[28]: finalSortedRDD = finalRDD.sortBy(lambda x: float(x[5]), True)
```

```
[ ]: # Add the header to RDD
```

```
[29]: headerList = [['station', 'station_longitude', 'station_latitude',
      ↪ 'day_of_week', 'hour', 'criticality_value']]
```

```
[30]: headerRDD = sc.parallelize(headerList)
```

```
[31]: csvFinal = headerRDD.union(finalSortedRDD)
```

```
[32]: def to_string(x):
      x[5] = str(x[5])
      return x
```

```
[33]: csvFinal = csvFinal.map(to_string)
```

```
[34]: finalTSVRDD = csvFinal.map(lambda x: '\t'.join(x))

[ ]: # save the result

[35]: finalTSVRDD.saveAsTextFile('critical-stations-Barcelona-RDD')

[36]: finalTSVRDD.collect()

[36]: ['station\tstation_longitude\tstation_latitude\tday_of_week\thour\tcriticality_value',
      '9\t2.185294\t41.385006\tFriday\t 10\t0.6129032258064516',
      '10\t2.185206\t41.384875\tSaturday\t 00\t0.622107969151671',
      '58\t2.170736\t41.377536\tMonday\t 01\t0.6239554317548747',
      '9\t2.185294\t41.385006\tFriday\t 22\t0.6258389261744967',
      '58\t2.170736\t41.377536\tMonday\t 00\t0.6323119777158774']
```

---

```
[ ]: # In this section, I am going to compute the distance between each station and
      ↳ the city center.
      # The city center has coordinates:
      # latitude = 41.386904
      # longitude = 2.169989
      # To compute the distance implement the Haversine function (use the formula
      # in https://en.wikipedia.org/wiki/Haversine_formula).
      # Then, compute the average number of used_slots per station

[ ]: # Define the function to compute the haversine

[66]: import math
      def haversine(x):
          lat1 = 41.386904
          lon1 = 2.169989
          # Radius of the Earth in kilometers
          R = 6371.0
          x[2] = float(x[2])
          x[1] = float(x[1])
          # Convert latitude and longitude from degrees to radians
          lat1, lon1, x[2], x[1] = map(math.radians, [lat1, lon1, x[2], x[1]])
          dlat = x[2] - lat1
          dlon = x[1] - lon1
          hav = math.sin(dlat / 2) ** 2 + math.cos(lat1) * math.cos(x[2]) * math.
          ↳ sin(dlon / 2) ** 2
          distance = 2 * R * math.asin(math.sqrt(hav))
          return x + [distance]
```

```

[ ]: # Add the distance to RDD

[43]: stationDistanceRDD = stationCleanRDD.map(haversine)

[ ]: # Create pair-RDD for (stations, distance)

[44]: stationDistancePair = stationDistanceRDD.map(lambda x: (x[0], x[4]))

[ ]: # Create pair-RDD for (stations, used_slots)

[45]: registerPair = registerCleanRDD.map(lambda l: (l[0], int(l[2])))

[ ]: # Group stations

[46]: registerPairGrouped = registerPair.groupByKey()

[ ]: # Calculate the average of used slots in each station

[47]: registerPairReduced = registerPairGrouped.map(lambda x: (x[0], sum(x[1])/
↪ len(x[1])))

[ ]: # Join station pair-RDD and register pair-RDD

[48]: distanceJoinedRDD = stationDistancePair.join(registerPairReduced)

[ ]: # Now, I want to find the stations that are closer than 1.5 km from the center

[ ]: # Filter distance closer than 1.5 km.

[49]: closeStations = distanceJoinedRDD.filter(lambda x: x[1][0] < 1.5)

[ ]: # calculate the number of stations closer than 1.5 km.

[50]: closeStations.count()

[50]: 64

[ ]: # Calculate the sum of used_slots of stations closer than 1.5 km to city center

[51]: closeStationsUsedSlots = closeStations.map(lambda x: float(x[1][1]))

[52]: closeStationsUsedSlotsSum = closeStationsUsedSlots.reduce(lambda a, b: a + b)

[53]: print(closeStationsUsedSlotsSum)

```

523.2437013187326

```
[ ]: # The average of used_slots of stations closer than 1.5 km from city center
```

```
[54]: avgCloseStationsUsedSlots = closeStationsUsedSlotsSum/closeStations.count()
```

```
[55]: print(avgCloseStationsUsedSlots)
```

8.175682833105197

```
[ ]: # Now, I am going to find the stations that are farther than 1.5 km from the
    ↪center
```

```
[ ]: # Filter distance further than 1.5 km.
```

```
[56]: furtherStations = distanceJoinedRDD.filter(lambda x: x[1][0] >= 1.5)
```

```
[ ]: # calculate the number of stations
```

```
[57]: furtherStations.count()
```

[57]: 220

```
[ ]: # Calculate the sum of used_slots of stations further than 1.5 km from city
    ↪center
```

```
[58]: furtherStationsUsedSlots = furtherStations.map(lambda x: float(x[1][1]))
```

```
[59]: furtherStationsUsedSlotsSum = furtherStationsUsedSlots.reduce(lambda a, b: a +
    ↪b)
```

```
[60]: print(furtherStationsUsedSlotsSum)
```

1731.242401514143

```
[ ]: # The average of used_slots of stations further than 1.5 km from city center
```

```
[61]: avgFurtherStationsUsedSlots = furtherStationsUsedSlotsSum/furtherStations.
    ↪count()
```

```
[62]: print(avgFurtherStationsUsedSlots)
```

7.869283643246105

```
[26]: # The result shows that the number of stations further than 1.5 km from city
    ↪center is approximately 3 times more
    # than those which are closer than 1.5 km from city center. Also, the average
    ↪of used slots of closer ones
```



# is higher than further ones. The average of used slots for closer stations is  $\hookrightarrow 8.17$  and it is 7.87 for further  
# ones that shows there are a little more free\_slots for further stations.